



Evening

Introduction

CS 571: *Operating Systems (Spring 2020)*
Lecture 1

Yue Cheng

Some material taken/derived from:

- Wisconsin CS-537 materials created by Remzi Arpacı-Dusseau.

Licensed for use under a Creative Commons Attribution-NonCommercial-ShareAlike 3.0 Unported License.

Introduction

- Instructor

- Dr. Yue Cheng (web: cs.gmu.edu/~yuecheng)
- Email: yuecheng@gmu.edu
- Office: 5324 Engineering
- Office hours: M 1:30pm-2:30pm
- Research interests: Distributed and storage systems, serverless and cloud computing, operating systems

HPC

Introduction

- Instructor
 - Dr. Yue Cheng (web: cs.gmu.edu/~yuecheng)
 - Email: yuecheng@gmu.edu
 - Office: 5324 Engineering
 - Office hours: M 1:30pm-2:30pm
 - Research interests: Distributed and storage systems, serverless and cloud computing, operating systems




- Teaching assistant
 - Abhishek Roy
 - Email: aroy6@masonlive.gmu.edu
 - Office hours:
 - TBD

Administrivia

OSTEP

- Required textbook
 - Operating Systems: Three Easy Pieces,
By Remzi H. Arpaci-Dusseau, Andrea C. Arpaci-Dusseau
- Recommended textbook
 - Operating Systems Principles & Practices
By T. Anderson and M. Dahlin
- **Prerequisites are enforced!!**
 - CS 310 Data Structures
 - CS 367 Computer Systems & Programming
 - CS 465 Computer Systems Architecture
 - Be comfortable with C programming language
- Class web page
 - <https://tddg.github.io/cs571-spring20/> *URL*
 - Class materials will all be available on the class web page



Administrivia (cont.)

- Syllabus
 - https://cs.gmu.edu/media/syllabi/Spring2020/CS_571ChengY.html
- Grading
 - 50% projects
 - 10% homework
 - 20% midterm exam
 - 20% final exam
- Reminders
 - Honor code
 - Late policy: 15% deducted each day. No credit after 3 days

Course schedule

- Materials, assignments, due dates

The screenshot shows the homepage of the CS 571 Operating Systems course at George Mason University. The page has a red header with the course name and university logo. Below the header, there's a navigation menu with links: Home, Course Information, Course Schedule, Project 0a (Linux utilities), Project 0b (Intro to OS/161), GitLab Setup, Installing OS/161, and Announcements. The 'Course Schedule' link is circled in blue. A blue curly brace on the right side groups the 'Project 0a' and 'Project 0b' links. A blue circle highlights the three-line menu icon in the top right corner.

The screenshot shows the 'Course Schedule' page for CS 571 Operating Systems (Spring 2020). The title 'Course Schedule' is underlined and circled in blue. A note below it states: 'The course schedule is tentative and subject to change.' A blue arrow points from the circled 'Course Schedule' on the left to this note. The schedule is presented in a table:

Week	Monday	Thursday
Week 1	Jan 20 MLK Day (NO CLASS)	Jan 23
Week 2	Jan 27 Lec 1: Introduction, process abstraction Proj 0a and Proj 0b out	Jan 30
Week 3	Feb 3 Lec 2: LDE, thread abstraction	Feb 6 Proj 0a and Proj 0b due Proj 1 out
Week 4	Feb 10 Lec 3: Synchronization I: locks, sem., and CV	Feb 13
Week 5	Feb 17 Lec 4: Synchronization II: classic sync problems, CPU scheduling I: FIFO, SJF	Feb 20 Proj 1a due
Week 6	Feb 24 (Traveling to FAST NO CLASS) Lec 5: CPU scheduling II: RR, priority, MLFQ	Feb 27
Week 7	Mar 2 Lec 6: Memory management I: address space, segmentation Midterm review Proj 2 out	Mar 5 Proj 1b due
Week 8	Mar 9 Spring recess (NO CLASS)	Mar 12 Enjoy or catchup?!

Course format

2 hr

10-15 min break



- (Review) + lecture + (worksheets and/or demos)
 - A short overview of the previous lecture to make sure the old content is not completely forgotten
 - Worksheet practices to make sure the lecture is well understood
 - Demos to help you gain a deeper understanding of the materials taught
 - OSTEP simulators, measurements

Course projects

- Goal:
 1. To gain hands-on systems programming experience
 2. To gain experience hacking a moderately sized system codebase (OS/161)

pthread fork

~20k Loc C ↗

Course projects

P0

P1

P2 {

P3

- Goal:

1. To gain hands-on systems programming experience
2. To gain experience hacking a moderately sized system codebase (OS/161)

MLFQ

Cmd-line interpreter

cat grep

zip / unzip

fork wait

lock / unlock

fork getpid

waitpid

exit ..

traffic light

thread mem sync

Course projects

- Goal:
 1. To gain hands-on systems programming experience
 2. To gain experience hacking a moderately sized system codebase (OS/161)
- Four coding projects (50%)
 - Project 0a (Warm-up): Linux utilities – 5% $\rightarrow 10\%$
 - Project 0b: Intro to OS/161 – 5% $\rightarrow 15\%$
 - Project 1a: Implement a Linux shell – 7% $\rightarrow 15\%$
 - Project 1b: OS/161 synchronization – 8% $\rightarrow 15\%$
 - Project 2a: OS/161 system calls – 10% $\rightarrow 15\%$
 - Project 2b: OS/161 CPU scheduling – 5%
 - Project 3: Implement a MapReduce app w/ C – 10%

Homework assignments

- Two written homework assignments
 - One before the midterm
 - One after the midterm

5% } 5% } 10%

Getting help

- Office hours
 - Monday 1:30 pm – 2:30 pm, Engineering 5324
 - Piazza
 - Good place to ask and answer questions
 - About project
 - About material from lecture
 - No anonymous posts or questions
- 


What is an OS?

What is an OS?

- OS manages resources
 - Memory, CPU, storage, network
 - Data (file systems, I/O)
- Provides low-level abstractions to applications
 - Files
 - Processes, threads
 - Virtual machines (VMs), containers
 - ...


H/W

OS abstracts away low-level details




OS abstracts away low-level details

- Under the surface
 - Complex and dirty implementations of abstractions and a lot more...



OS abstracts away low-level details

- User's perspective
 - User interface:
 - Terminal, GUI
 - Application interface:
 - System calls *API vs*
- Under the surface
 - Complex and dirty implementations of abstractions and a lot more...



The goals of an OS

- OS manages resources
 - Memory, CPU, storage, network
 - Data (file systems, I/O)
- Provides low-level abstractions to applications
 - Files
 - Processes, threads
 - Virtual machines (VMs), containers
 - ...
- Goals
 - Resource efficiency (resource virtualization)
 - Ease-of-use (interfaces)
 - Reliability (user-kernel space separation)

std C


malloc

syscalls

System Calls

- System calls provide the interface between a running program and the operating system
 - Generally available in routines written in C and C++
 - Certain low-level tasks may have to be written using assembly language
- Typically, application programmers design programs using an application programming interface (API)
- The runtime support system (runtime libraries) provides a system-call interface, that intercepts function calls in the API and invokes the necessary system call within the operating system
- Major differences in how they are implemented (e.g., Windows vs. Unix)

Example System Call Processing

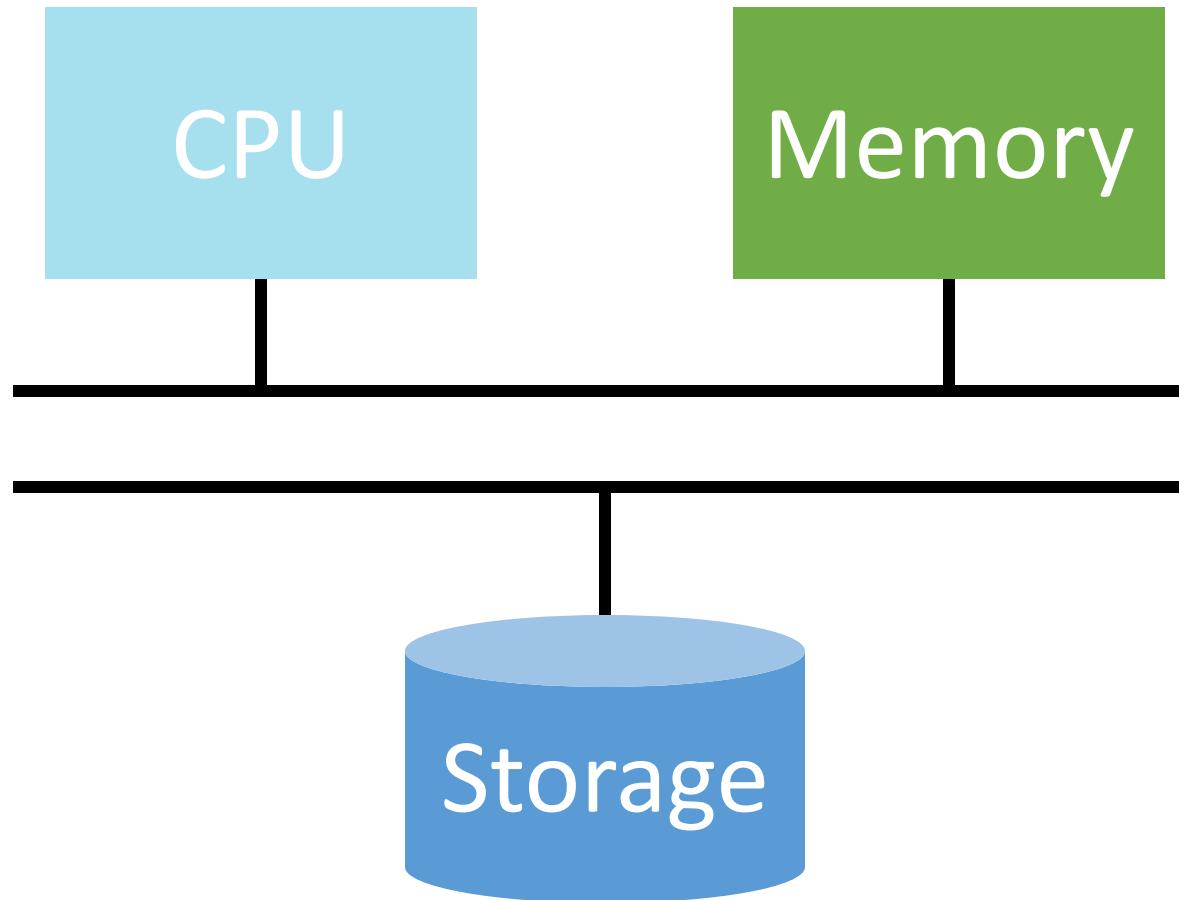


Major System Calls in Linux: File Management

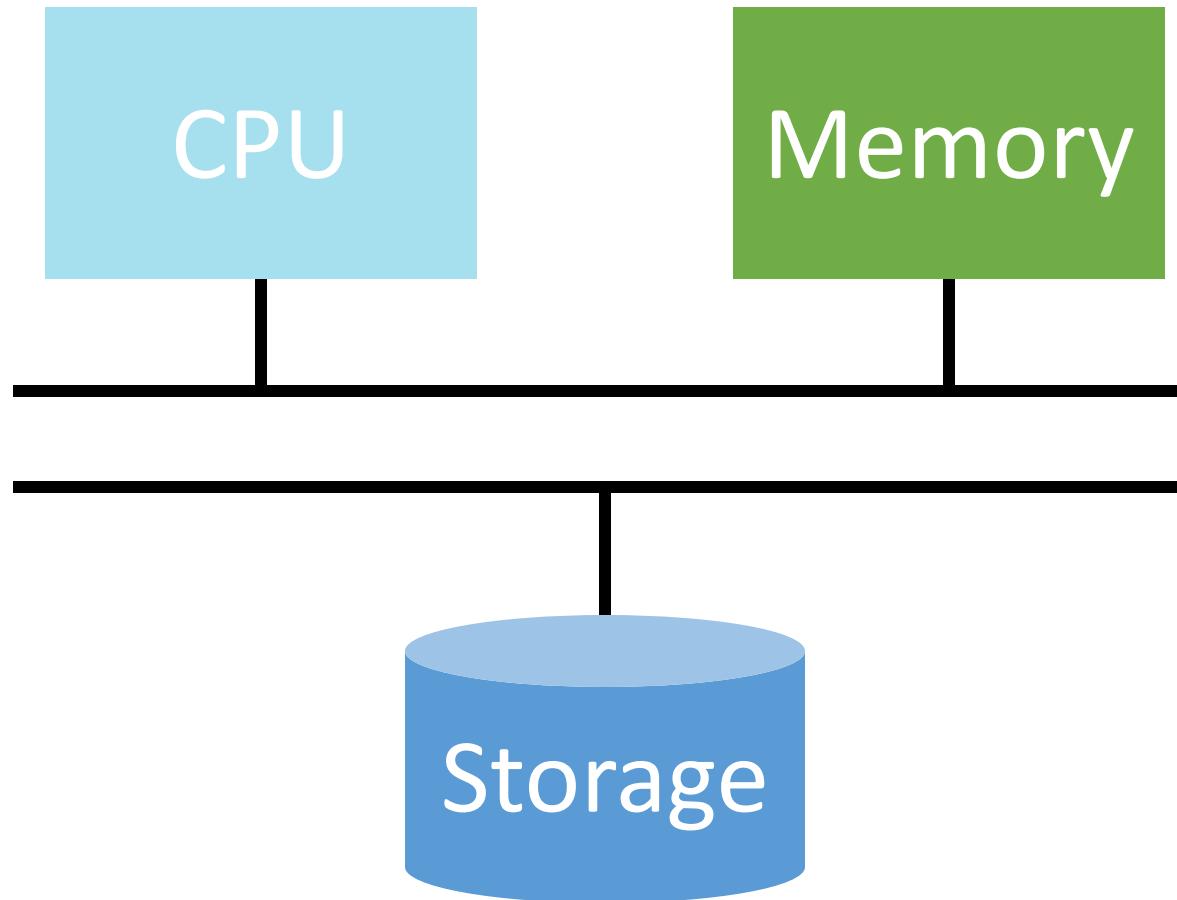
int

- fd = open(file, how, ...)
 - Open a file for reading, writing, or both
- s = close(file) → fd
 - Close an open file
- n = read(fd, buf, nbytes)
 - Read data from a file into a buffer
- n = write(fd, buf, nbytes)
 - Write data from a buffer into a file
- pos = lseek(fd, offset, whence)
 - Move the file pointer
- s = stat(name, &buf)
 - Get a file's status info

3 Major Topics



OS Provides Virtualization on Hardware



Topic 1: Concurrency, Synchronization, and CPU Scheduling


- Process/thread abstraction
- Synchronization
- CPU scheduling



Process Abstraction

- A process is a program in execution
 - It is a unit of work within the system. A program is a **passive entity**, a process is an **active entity**.
- Process needs resources to accomplish its task
 - CPU, memory, I/O, files
 - Initialization data
- Process termination requires reclaim of any reusable resources
- Single-threaded process has one program counter specifying location of next instruction to execute
 - Process executes instructions sequentially, one at a time, until completion
- Multi-threaded process has one program counter per thread
- A software system may have many processes, some user, some operating system running concurrently on one or more CPUs
 - Concurrency by multiplexing the CPUs among the processes / threads

Loading from Program to Process



Topic 2: Memory Management and Virtual Memory

- Process/thread abstraction
 - Synchronization
 - CPU scheduling
- Memory management
 - Virtual memory

time-shared

space-shared


LRU

Storage

Memory Management

- All data in memory before and after processing
- All instructions in memory in order to execute
- Memory management determines what is in memory when
 - Optimizing CPU utilization and computer response to users
- Memory management activities PT
 - Keeping track of which parts of memory are currently being used and by whom
 - Deciding which processes (or parts thereof) and data to move into and out of memory
 - Allocating and deallocating memory space as needed
- **Virtual memory** management is an essential part of most operating systems

Topic 3: Storage, I/O, and Filesystems




Storage Management

- OS provides a uniform, logical view of information storage
 - Abstracts physical properties to logical storage unit - file
 - Each medium is controlled by device type (i.e., disk drive, tape drive)
 - Varying properties include access speed, capacity, data-transfer rate, access method (sequential or random)
- Filesystem management
 - Files usually organized into directories
 - Access control on most systems to determine who can access what
 - OS activities include
 - Creating and deleting files and directories
 - Primitives to manipulate files and dirs
 - Mapping files onto secondary storage
 - Backup files onto stable (non-volatile) storage media

Storage hierarchy

tradeoff

Fastest
\$ \$ \$ \$



Storage Structure

- **Main memory** – relatively large storage media that the CPU can access directly
 - Small CPU cache memories are used to speed up average access time to the main memory at run-time
 - Volatile (data loss at power-off)
 - Byte-addressable
- **Secondary storage** – extension of main memory that provides large nonvolatile storage capacity.
 - Magnetic disks
 - Electronic disks -- Solid state disks (SSDs)
 - Non-volatile (i.e., persistent)
 - Non byte-addressable

Storage Systems Tradeoffs

- Storage systems organized in hierarchy
 - Speed
 - Cost
 - Volatility
 - Density
- Faster access time, greater cost per bit
- Greater capacity (density), lower cost per bit
- Greater capacity (density), slower access speed

Increased complexity – Memory

2015



Byte-addressable

L1/L2 cache

~1 ns

} CPU

L3 cache

~10 ns

Main memory

DRAM BW capacity

~100 ns / ~80 GB/s / ~100GB

Lat

NAND SSD

~100 usec / ~10 GB/s / ~1 TB

Fast HDD

~10 msec / ~100 MB/s / ~10 TB

4KB
512B

Increased complexity – Memory

2015



L1/L2 cache

~1 ns

L3 cache

~10 ns

Main memory

~100 ns / ~80 GB/s / ~100GB

NAND SSD

~100 usec / ~10 GB/s / ~1 TB

Fast HDD

~10 msec / ~100 MB/s / ~10 TB

2020



L1/L2 cache

~1 ns

L3 cache

~10 ns

HBM

Main memory

~10 ns

NVM (Intel Optane DC)

NAND SSD

~10 ns / ~1TB/s / ~10GB

Main memory

~100 ns / ~80 GB/s / ~100GB

NVM (Intel Optane DC)

~1 usec / ~10GB/s / ~1TB

NAND SSD


~100 usec / ~10 GB/s / ~10 TB

Fast HDD

~10 msec / ~100 MB/s / ~100 TB

The CPU-Memory Gap

The gap widens between memory, disk, and CPU speeds.




Caching

80-20


- Skew rule: 80% requests hit on 20% hottest data
- Important principle, performed at many levels in a computer (in hardware, operating system, software)
- Information in use copied from slower to faster storage temporarily
- Faster storage (cache) checked first to determine if information is there
 - If it is, information used directly from the cache (fast)
 - If not, data copied to cache and used there
- Cache smaller than storage being cached
 - Cache management important design problem
 - Cache size and replacement policy

Intel Core i7 Cache Hierarchy



Migration of Integer A from Disk to Register

- Multitasking environments must be careful to use most recent value, no matter where it is stored in the storage hierarchy



- Multiprocessor environment must provide **cache coherency** in hardware such that all CPUs have the most recent value in their cache
- Distributed environment situation even more complex
 - Several copies of a piece of data can exist

Why do you take this course?

General Learning Goals

1. Grasp **basic** knowledge about **Operating Systems** and **Computer Systems** software
2. Learn **important systems concepts** in general
 - Multi-processing/threading, synchronization
 - Scheduling
 - Caching, memory, storage
 - And more...
3. Gain **hands-on** experience in **writing/hacking/designing** moderately large systems software

Why do you take this course?

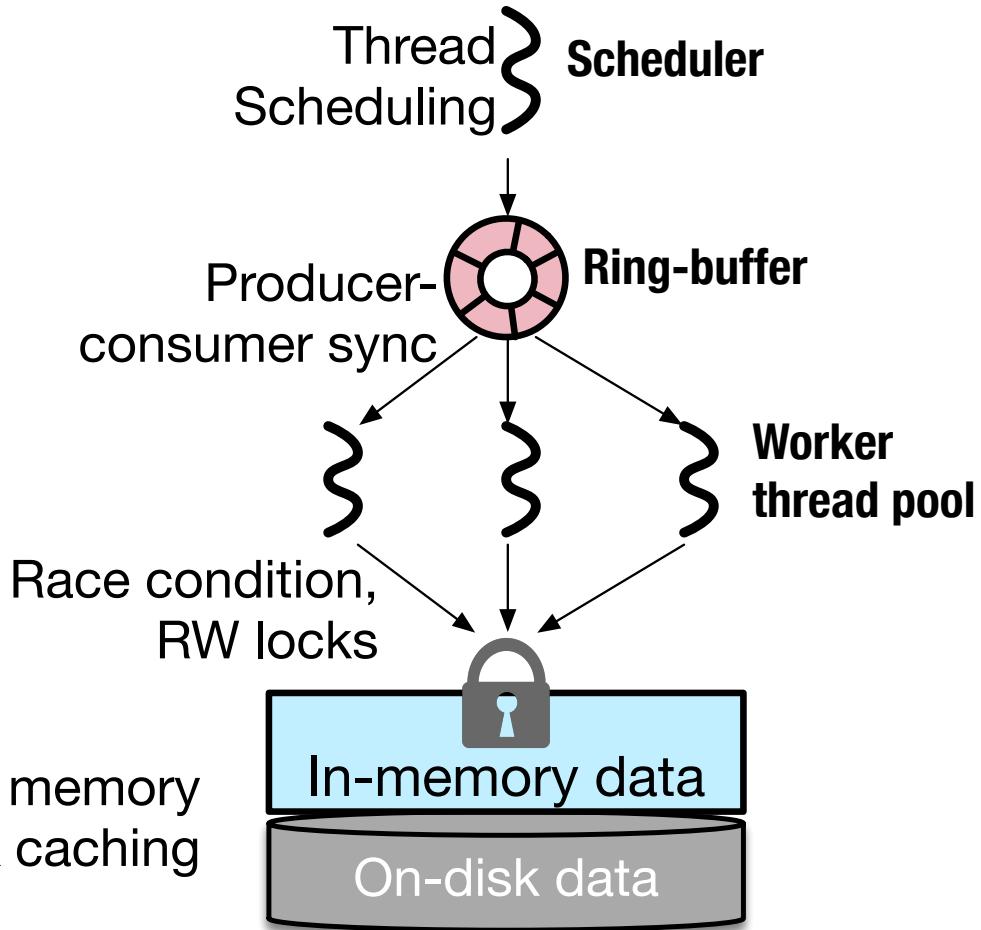
- The OS concepts are everywhere
 - Fundamental OS techniques broadly generalize to widely-used systems technique
 - Scheduling
 - Concurrency
 - Memory management
 - Caching
 - ...

One example: Memcached



- Memcached is a distributed in-memory object cache system
 - Written in C
 - In-memory hash table
 - Multi-threading

Virtual memory & caching



Memcached can be treated as a user-space mini-OS

What is a Process?

What is a Process?

- **Programs** are code (static entity)
- **Processes** are running programs
- Java analogy
 - class -> “program”
 - object -> “process”


What is in a Process?

Process




What things change as a program runs?

What is in a Process?



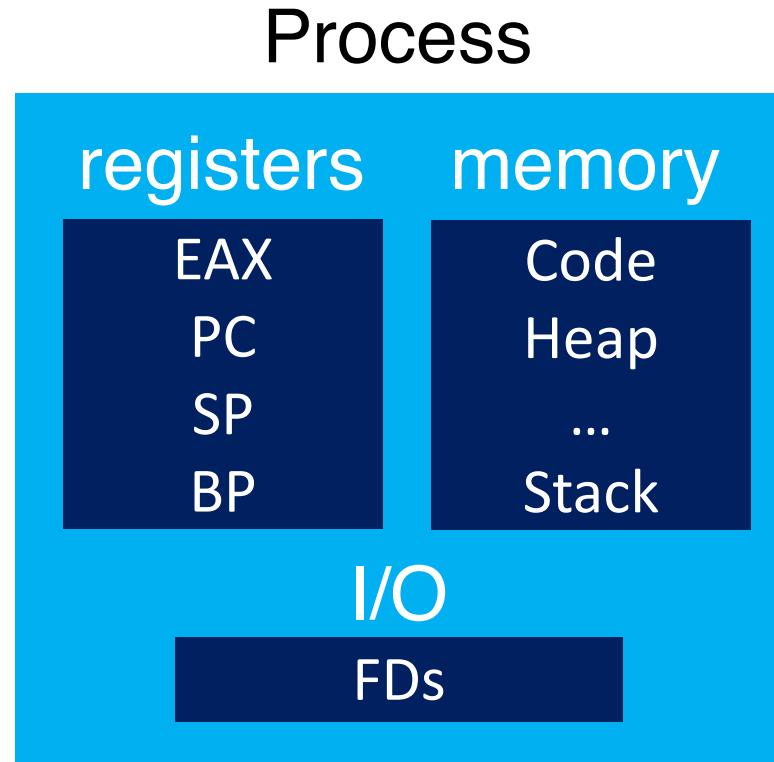
What things change as a program runs?

What is in a Process?



What things change as a program runs?

What is in a Process?



What things change as a program runs?


Peeking Inside

- Processes share code, but each has its own “**context**”
- CPU
 - Instruction pointer (Program Counter)
 - Stack pointer
- Memory
 - Set of memory addresses (“**address space**”)
 - cat /proc/<PID>/maps
- Disk
 - Set of file descriptors
 - cat /proc/<PID>/fdinfo/*

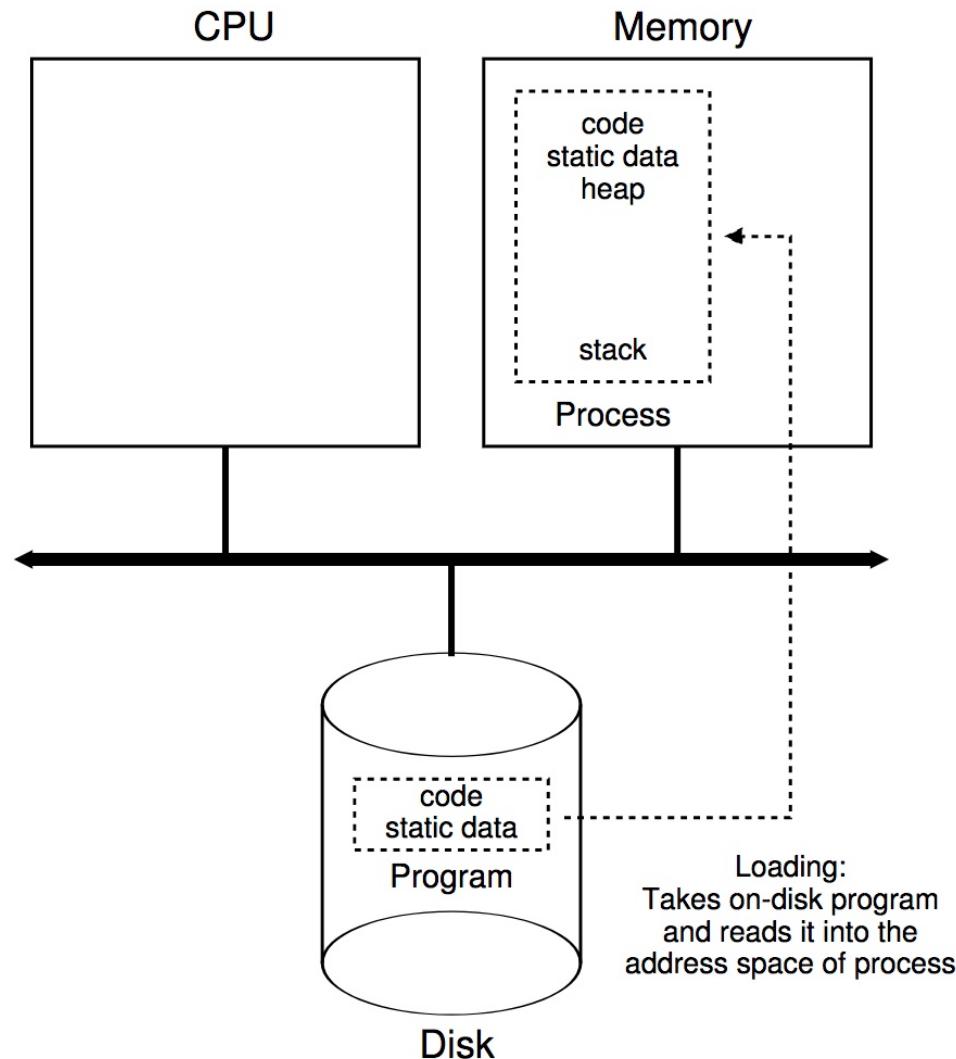
Process Creation

- Principle events that cause process creation
 - System initialization
 - Execution of a process creation system call by a running process
 - User request to create a process


Process Creation



Process Creation



Process Creation




Process Creation (cont.)

- Parent process creates children processes, which, in turn create other processes, forming a tree (**hierarchy**) of processes
- **Questions:**
 - Will the parent and child execute **concurrently**?
 - How will the **address space** of the child be related to that of the parent?
 - Will the parent and child **share some resources**?

An Example Process Tree


pid = 0



How to View Process Tree in Linux?

- % `ps auxf`
 - ‘f’ is the option to show the process tree
- % `pstree`

Process Creation in Linux



- Each process has a **process identifier (pid)**
- The parent executes **fork()** system call to spawn a child
- The child process has a **separate copy** of the parent's address space
 - Both the parent and the child continue execution at the instruction following the **fork()** system call
 - The return value for the **fork()** system call is
 - **zero** value for the new (**child**) process
 - **non-zero pid** for the **parent** process
 - Typically, a process can execute a system call like **exec1()** to load a binary file into memory

Process Creation in Linux

- Each process has a **process identifier (pid)**
- The parent executes **fork()** system call to spawn a child
- The child process has a **separate copy** of the parent's address space
 - Both the parent and the child continue execution at the instruction following the **fork()** system call
 - The return value for the **fork()** system call is
 - **zero** value for the new (**child**) process
 - **non-zero pid** for the **parent** process
 - Typically, a process can execute a system call like **exec()** to load a binary file into memory

This is really the pid of the child process

The man page of fork()

<http://man7.org/linux/man-pages/man2/fork.2.html>

RETURN VALUE

[top](#)

On success, the PID of the child process is returned in the parent, and 0 is returned in the child. On failure, -1 is returned in the parent, no child process is created, and `errno` is set appropriately.

ERRORS

[top](#)

EAGAIN A system-imposed limit on the number of threads was encountered. There are a number of limits that may trigger this error:

- * the **RLIMIT_NPROC** soft resource limit (set via `setrlimit(2)`), which limits the number of processes and threads for a real user ID, was reached;
- * the kernel's system-wide limit on the number of processes and threads, `/proc/sys/kernel/threads-max`, was reached (see `proc(5)`);
- * the maximum number of PIDs, `/proc/sys/kernel/pid_max`, was reached (see `proc(5)`); or
- * the PID limit (**pids.max**) imposed by the cgroup "process number" (PIDs) controller was reached.

Example Program with fork()

```
void main () {
    int pid;

    pid = fork();
    if (pid < 0) {/* error_msg */}
    else if (pid == 0) { /* child process */
        execl("/bin/ls", "ls", NULL); /* execute ls */
    } else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        exit(0);
    }
    return;
}
```

A Very Simple Shell using fork()

```
while (1) {  
    type_prompt();  
    read_command(cmd);  
    pid = fork();  
    if (pid < 0) {/* error_msg */}  
    else if (pid == 0) { /* child process */  
        execute_command(cmd);  
    } else { /* parent process */  
        wait(NULL);  
    }  
}
```

More example: fork 1

```
forkexample.c *  
1 #include <sys/types.h>  
2 #include <stdio.h>  
3 #include <stdlib.h>  
4 #include <unistd.h>  
5  
6 int number = 7;  
7  
8 int main(void) {  
9     pid_t pid;  
10    printf("\nRunning the fork example\n");  
11    printf("The initial value of number is %d\n", number);  
12  
13    pid = fork();  
14    printf("PID is %d\n", pid);  
15  
16    if (pid == 0) {  
17        number *= number;  
18        printf("\tIn the child, the number is %d -- PID is %d\n", number, pid);  
19        return 0;  
20    } else if (pid > 0) {  
21        wait(NULL);  
22        printf("In the parent, the number is %d\n", number);  
23    }  
24  
25    return 0;  
26}  
27
```

Results

./forkexample1

Running the fork example

The initial value of number is 7

PID is 2137

PID is 0

In the child, the number is 49 -- PID is 0

In the parent, the number is 7

Further more example: fork 2

```
forkexample2.c *  
1 #include <sys/types.h>  
2 #include <stdio.h>  
3 #include <stdlib.h>  
4 #include <unistd.h>  
5  
6 int number = 7;  
7  
8 int main(void) {  
9     pid_t pid;  
10    printf("\nRunning the fork example\n");  
11    printf("The initial value of number is %d\n", number);  
12  
13    pid = fork();  
14    printf("PID is %d\n", pid);  
15  
16    if (pid == 0) {  
17        number *= number;  
18        fork();  
19        printf("\tIn the child, the number is %d -- PID is %d\n", number, pid);  
20        return 0;  
21    } else if (pid > 0) {  
22        wait(NULL);  
23        printf("In the parent, the number is %d\n", number);  
24    }  
25  
26    return 0;  
27 }  
28 }
```

Results

./forkexample2

Running the fork example

The initial value of number is 7

PID is 2164

PID is 0

In the child, the number is 49 -- PID is 0

In the child, the number is 49 -- PID is 0

In the parent, the number is 7

execl (or execvp) vs. fork

```
execlexample.c *  
1 #include <sys/types.h>  
2 #include <stdio.h>  
3 #include <stdlib.h>  
4 #include <unistd.h>  
5  
6 int number = 7;  
7  
8 int main(void) {  
9     pid_t pid;  
10    printf("\nRunning the execl example\n");  
11    pid = fork();  
12    printf("PID is %d\n", pid);  
13  
14    if (pid == 0) {  
15        printf("\tIn the execl child, PID is %d\n", pid);  
16        execl("./forkexample2", "forkexample2", NULL);  
17        return 0;  
18    } else if (pid > 0) {  
19        wait(NULL);  
20        printf("In the parent, done waiting\n");  
21    }  
22  
23    return 0;  
24 }
```

Results

./execlexample

Running the execl example
PID is 2179
PID is 0

In the execl child, PID is 0

Running the fork example
The initial value of number is 7
PID is 2180
PID is 0

In the child, the number is 49 -- PID is 0

In the child, the number is 49 -- PID is 0


In the parent, the number is 7
In the parent, done waiting

} forkexample2


Today's demo code

- You can fork it here:
<https://github.com/tddg/demo-ostep-code>
 - under `cpu-api/`

Process Creation




Process Creation



Now, after process creation, CPU begins directly executing process code

Process Creation



Challenge: how to prevent process from doing “OS kernel stuff”?